


Programming in R



Unit 9: Randomness



Reproducibility

Remember, this course has multiple goals:

- Learn things about the R language: "R"
- Get to know nice tools to use: "Tools"
- Learn things about software development in general: "Dev"

This unit:

- "R" Track: Random Numbers

R Track

Random Numbers

Random Numbers

There are many reasons why your code could want to have **randomness**

- Sample points from a probability distribution for a simulation
- Probabilistic algorithms, e.g. for estimating the expected value of something by randomly evaluating points and taking their average
- Simulate dice or shuffled cards for a game
- Create a password or key for encryption

Random Numbers

This is fundamentally at odds with our concept of computers and how we program them: programs are *deterministic*.

Solutions

- Use [hardware-generated randomness](#), such as electrical noise in an analog circuit -- slow and expensive!
- Use a [pseudorandom number generator](#) (PRNG, sometimes just called RNG): produce numbers that are calculated deterministically, but that have a pattern so complicated that they are essentially random for their purpose -- fast, but not "really" "random"
 - Like using the digits of pi, or digits of sqrt(2): deterministic, but not correlated with most things one would be doing with it
 - Actual algorithms are documented / referenced in `?Random`.
- Pseudorandom numbers in R: generated with functions starting with 'r' (runif, rexp, rbinom, etc. -- see `?Distributions`), as well as 'sample' / 'sample.int'; some others.


Example PRNG: "Middle-Square Method"

1. Take a number "x" with n digits (e.g. n = 2)

x = 43

Example PRNG: "Middle-Square Method"


1. Take a number "x" with n digits (e.g. n = 2)
2. set $x \rightarrow x^2$



Example PRNG: "Middle-Square Method"


1. Take a number "x" with n digits (e.g. n = 2)
2. set $x \rightarrow x^2$

Restrict x to the middle n digits




Example PRNG: "Middle-Square Method"

1. Take a number "x" with n digits (e.g. $n = 2$)
2. set $x \rightarrow x^2$
3. Restrict x to the middle n digits
4. Repeat from 2.



Example PRNG: "Middle-Square Method"

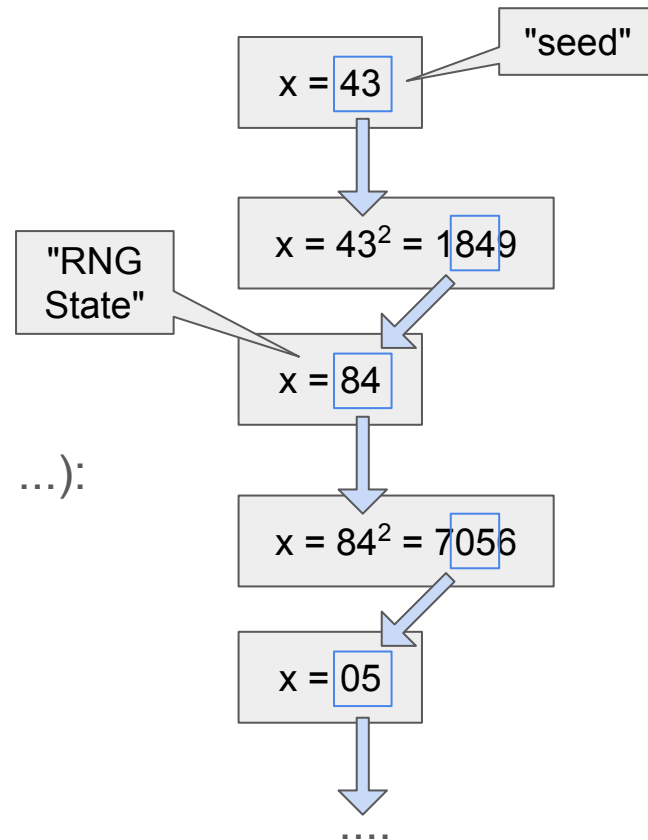
1. Take a number "x" with n digits (e.g. $n = 2$)
 2. set $x \rightarrow x^2$
 3. Restrict x to the middle n digits
 4. Repeat from 2.
- Initial x is our "**seed**"



Example PRNG: "Middle-Square Method"

1. Take a number "x" with n digits (e.g. n = 2)
2. set $x \rightarrow x^2$
3. Restrict x to the middle n digits
4. Repeat from 2.


- Initial x is our "**seed**"
- We get a series of x values (43, 84, 05, 25, 62, ...):
 "**RNG State**"



Example PRNG: "Middle-Square Method"

1. Take a number "x" with n digits (e.g. n = 2)
2. set $x \rightarrow x^2$
3. Restrict x to the middle n digits
4. Repeat from 2.

- Initial x is our "**seed**"
- We get a series of x values (43, 84, 05, 25, 62, ...):
 "**RNG State**"
- We could simulate a coin toss: is the state even or odd?
 → H, T, H, H, T, T,

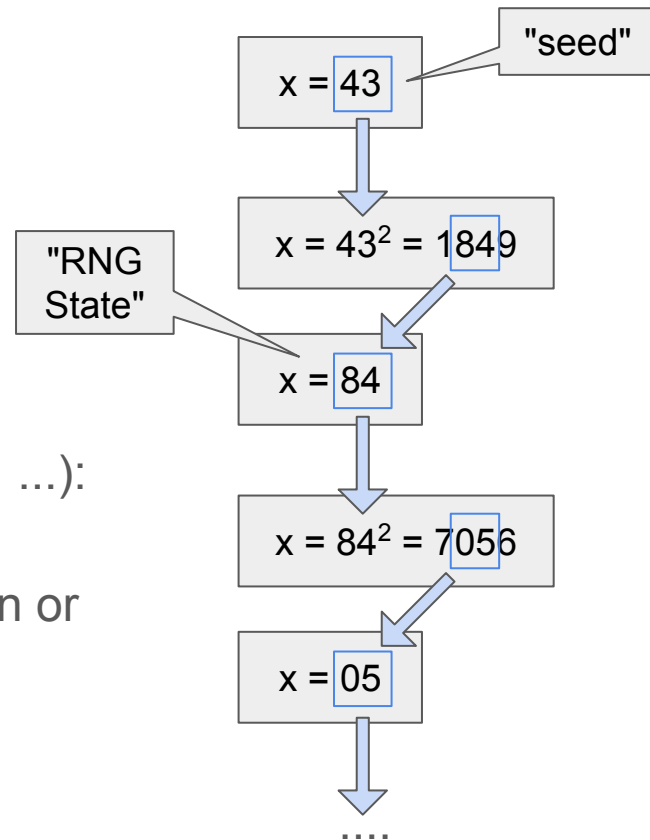


Example PRNG: "Middle-Square Method"

1. Take a number "x" with n digits (e.g. n = 2)
2. set $x \rightarrow x^2$
3. Restrict x to the middle n digits
4. Repeat from 2.

- Initial x is our "**seed**"
- We get a series of x values (43, 84, 05, 25, 62, ...):
 "**RNG State**"
- We could simulate a coin toss: is the state even or odd?

→ H, T, H, H, T, T,
 (43, 84, 05, 25, 62, 84, ...)



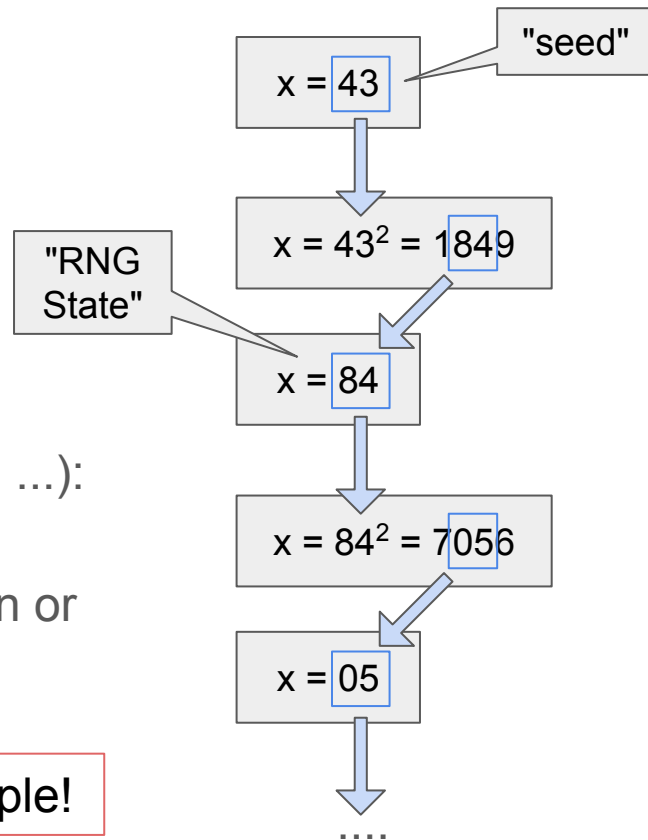
Example PRNG: "Middle-Square Method"

1. Take a number "x" with n digits (e.g. n = 2)
2. set $x \rightarrow x^2$
3. Restrict x to the middle n digits
4. Repeat from 2.

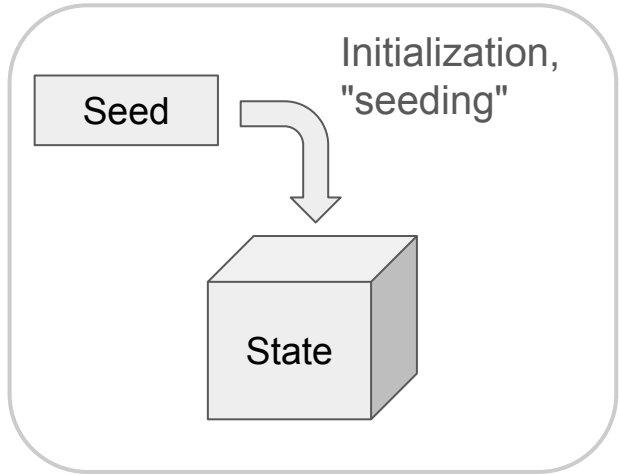
- Initial x is our "**seed**"
- We get a series of x values (43, 84, 05, 25, 62, ...):
 "**RNG State**"
- We could simulate a coin toss: is the state even or odd?

→ H, T, H, H, T, T,
 (43, 84, 05, 25, 62, 84, ...)

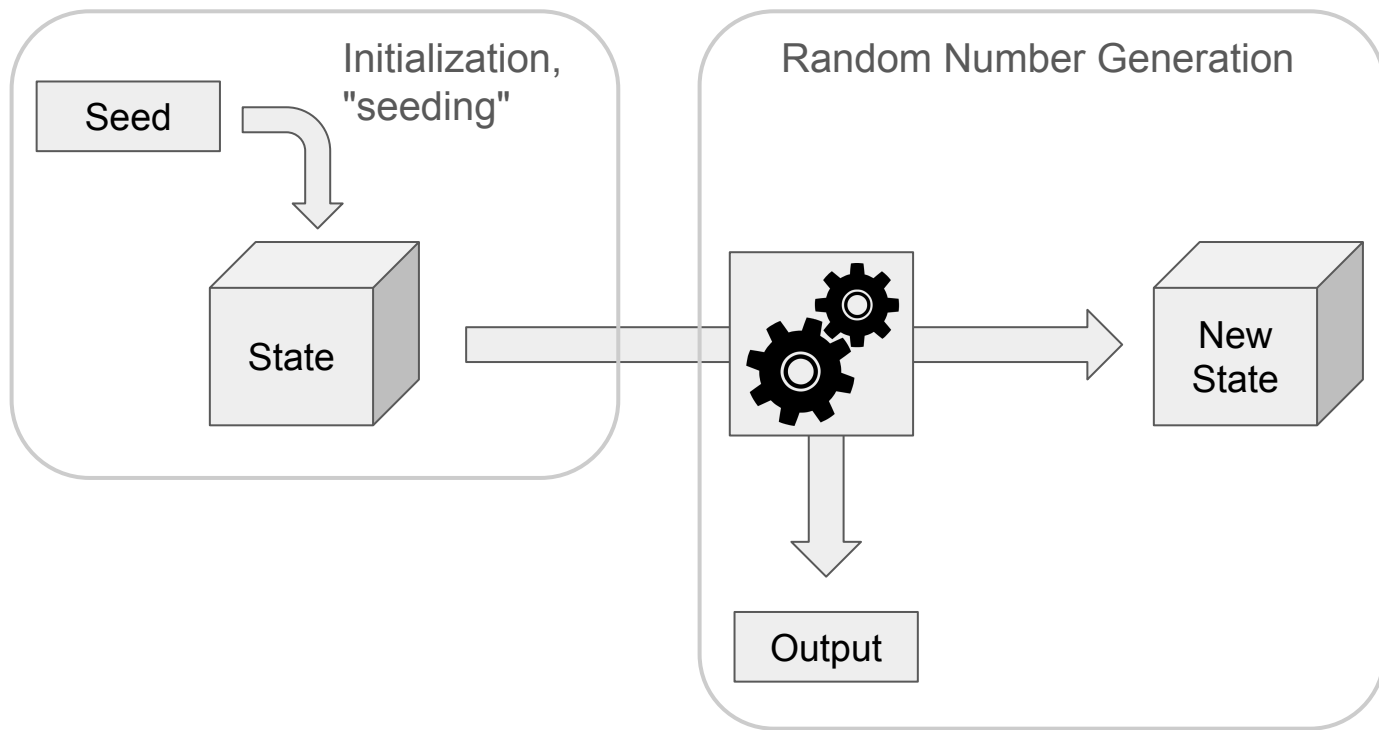
This is a toy-example!




PRNG Generally



PRNG Generally



PRNG in R: Seeding




`.Random.seed`

- see it with `ls(all.names = TRUE)`
- only exists once you `set.seed()` or use a random function
- A large vector of numbers!

```
> .Random.seed
```

```
[1]          403          287 -1740373036    152930328 -1474854045 -1128417688
[7] 1080437759 1475710325 -2031914562    374687669 -1186119422    385312014
[13]  797075817 1614705622  1012171569  1072024031   457256672 -519501961
[19]  288275904  329852028   119518183   220589868  2100395579 -1738058023
[25] -813976942 -2067108231 -1365388634 -2026736262 -192084811  -210639310
[31] -1189681291 -644081813   144763212  1994894451 -2026972484 -1842039744
```

PRNG in R: Seeding



`.Random.seed`

- see it with `ls(all.names = TRUE)`
- only exists once you `set.seed()` or use a random function
- A large vector of numbers!
- When using the same RNG algorithm (`RNGkind()`),
the same seed will lead to the same initial state

• `set.seed(1)`


• `.Random.seed`

Probably
the same
for you!

```
[1] 10403      624 -169270483 -442010614 -603558397 -222347416
1489374793 865871222 1734802815 98005428 268448037 63650722
-1754793285 -2135275840 -779982911 -864886130 1880007095 463784588
1271615005 1390544442 -544608653 -251475688 -326549447 -1570483546
1965989103 -784675228 1458985493 2146317266 -1103943381 289023600
[31] 436963407 109630910 69979943 1606475068 1441346829 -662821782
```

PRNG in R: Drawing

- Functions starting with `r`:
 - `runif()`: uniform distribution
 - `rexp()`: exponential distribution
 - `rnorm()`: normal distribution
 - ...




PRNG in R: Drawing

- Functions starting with `r`:
 - `runif()`: uniform distribution
 - `rexp()`: exponential distribution
 - `rnorm()`: normal distribution
 - ...
- Internally, when drawing random value:
 - draw one or more uniformly distributed values
 - build other random variables from this.

One could, e.g., do:

```
rexp() = -log(runif())
```


(R does something like this but more complicated)



PRNG in R: Drawing


Drawing RVs advances the state:

```
> set.seed(1)
> head(.Random.seed)
[1] 10403      624 -169270483 -442010614 -603558397 -222347416
> runif(1)
[1] 0.2655087
> head(.Random.seed)
[1] 10403      1 1654269195 -1877109783 -961256264 1403523942
```



PRNG in R: Drawing

Drawing RVs advances the state:



```
> set.seed(1)
> head(.Random.seed)
[1] 10403      624 -169270483 -442010614 -603558397 -222347416
> runif(1)
[1] 0.2655087
> head(.Random.seed)
[1] 10403      1 1654269195 -1877109783 -961256264 1403523942
```

This is deterministic! Do it again, it advances the same amount.

```
> set.seed(1)
> head(.Random.seed)
[1] 10403      624 -169270483 -442010614 -603558397 -222347416
> runif(1)
[1] 0.2655087
> head(.Random.seed)
[1] 10403      1 1654269195 -1877109783 -961256264 1403523942
```

PRNG in R: Drawing

Important for reproducibility:

- Setting the same seed gives the same random values.
- Calling the same sequence of random functions in the same order(!) yields the same sequence of results
- The RNG state advances in discrete steps

Reset

```
> set.seed(1)
> runif(2)
[1] 0.2655087 0.3721239
```

Reset

```
> runif(2)
[1] 0.5728534 0.9082078
> set.seed(1)
> runif(2)
[1] 0.2655087 0.3721239
```


PRNG in R: Drawing

Important for reproducibility:

- Setting the same seed gives the same random values.
- Calling the same sequence of random functions in the same order(!) yields the same sequence of results
- The RNG state advances in discrete steps

Reset

```
> set.seed(1)
> runif(2)
[1] 0.2655087 0.3721239
```

Reset

```
> runif(2)
[1] 0.5728534 0.9082078
```

Reset

```
> set.seed(1)
> runif(2)
[1] 0.2655087 0.3721239
```

```
> set.seed(1)
> runif(1)
[1] 0.2655087
> runif(1)
[1] 0.3721239
```

PRNG in R: Drawing

Important for reproducibility:

- Setting the same seed gives the same random values.
- Calling the same sequence of random functions in the same order(!) yields the same sequence of results
- The RNG state advances in discrete steps
- Some calls advance the state more than others

```
> set.seed(1) Reset
> runif(6)
[1] 0.2655087 0.3721239 0.5728534 0.9082078 0.2016819 0.8983897
> set.seed(1) Reset
> rnorm(1)
[1] -0.6264538
> runif(6)
[1] 0.5728534 0.9082078 0.2016819 0.8983897 0.9446753 0.6607978
```

PRNG in R: Drawing

Important for reproducibility:

- Setting the same seed gives the same random values.
- Calling the same sequence of random functions in the same order(!) yields the same sequence of results
- The RNG state advances in discrete steps
- Some calls advance the state more than others

```
> set.seed(1) Reset
> runif(6)
[1] 0.2655087 0.3721239 0.5728534 0.9082078 0.2016819 0.8983897
> set.seed(1) Reset
> rnorm(1)
[1] -0.6264538
> runif(6)
[1] 0.5728534 0.9082078 0.2016819 0.8983897 0.9446753 0.6607978
> set.seed(1) Reset
> rnorm(2)
[1] -0.6264538 0.1836433
> runif(6)
[1] 0.2016819 0.8983896 0.9446752 0.6607977 0.6291140 0.0617862
```

PRNG in R: Drawing

Important for reproducibility:

- Setting the same seed gives the same random values.
- Calling the same sequence of random functions in the same order(!) yields the same sequence of results
- The RNG state advances in discrete steps
- Some calls advance the state more than others
- `.Random.seed` saves the state
- but beware that some random functions may "cache" their values & have addnl. state -- see `?Random`
--> use `set.seed()`!

```
> set.seed(1)
> runif(1)
[1] 0.2655087
> x <- .Random.seed
> runif(1)
[1] 0.3721239
> .Random.seed <- x
> runif(1)
[1] 0.3721239
```

Using PRNGs

Usecase: Hypothesis testing & p-value estimation

- You measure a statistic that is transformed in a complicated way?
 - e.g. rounding, cutoffs, ...
 - -> Simulate the data under H_0 to calculate p-value!
- *Bootstrapping*: draw samples with replacement and fit models
 - `sample(values, replace = TRUE)`
- Permutation tests
 - Is data column X relevant for a given result?
 - Shuffle this column and run the analysis again (x1000) and see what changes
 - (Does not always work: E.g. if X and Y are almost copies of each other, shuffling only one will sometimes have no visible effect)

Usecase: Random Search

- You have something with "hyperparameters" -- settings that influence performance
E.g. $f(x, y, z)$
- Often some of them are important, others have no influence
E.g. x, y have an influence, z does not
In general, you don't know which!
- If you try a grid of 10 values each, you need to evaluate $f()$ $10 \times 10 \times 10$ times (in general: 10^N)
- If you randomly sample values, you get a good resolution in $\sim 10 \times 10$ evals!

[[Bergstra & Bengio 2012](#)]

Usecase: Visualization & Data Exploration

- If you have a very large dataset
- You want to explore it in some way:
 - plot it along various axes
 - try out various methods
- You often get a good approximation of what it looks like / what works well if you consider a small subset
- -> Randomly subsample your data

Usecase: More Specific Algorithms

Many algorithms make use of randomness more specifically:

- Markov-Chain Monte-Carlo -- Sample from & integrate complex probability distributions, relevant in Bayesian modelling
- Evolutionary algorithms: Quite robust methods for optimization
- Data augmentation techniques

These are often implemented in packages, you don't need to write these yourself.

Distributions in R

	Random Value	Probability (density)	Cumulative d.f., $P(X \leq x)$	Quantile function, F^{-1}
Uniform dist.	<code>runif()</code>	<code>dunif()</code>	<code>punif()</code>	<code>qunif()</code>
Normal dist.	<code>rnorm()</code>	<code>dnorm()</code>	<code>pnorm()</code>	<code>qnorm()</code>
Exponential	<code>rexp()</code>
Binomial	<code>rbinom()</code>	...		
Geometric	<code>rgeom()</code>	...		
Poisson	<code>rpois()</code>	...		
...	...			

Random Numbers in R: Recipes

Do a random experiment 100 times

Random Numbers in R: Recipes

Do a random experiment 100 times

Dont:

```
result <- numeric(n)
for (i in seq_len(n)) {
  result[[i]] <- experiment()
}
```

Random Numbers in R: Recipes

Do a random experiment 100 times

Dont:

```
result <- numeric(n)
for (i in seq_len(n)) {
  result[[i]] <- experiment()
}
```

or:

```
result <- vapply(seq_len(n), function(i) {
  experiment()
}, numeric(1))
```

Random Numbers in R: Recipes

Do a random experiment 100 times

Dont:

```
result <- numeric(n)
for (i in seq_len(n)) {
  result[[i]] <- experiment()
}
```

or:

```
result <- vapply(seq_len(n), function(i) {
  experiment()
}, numeric(1))
```

Do:

```
result <- replicate(
  n,
  experiment()
)
```

Random Numbers in R: Recipes

Do a random experiment 100 times

Dont:

```
result <- numeric(n)
for (i in seq_len(n)) {
  result[[i]] <- experiment()
}
```



or:

```
result <- vapply(seq_len(n), function(i) {
  experiment()
}, numeric(1))
```



Do:

```
result <- replicate(
  n,
  experiment()
)
```



Use `replicate()` if
your expression does
not depend on `i`!

Random Numbers in R: Recipes

Get a random integer between 1 and `n`

Random Numbers in R: Recipes

Get a random integer between 1 and n

Dont:

```
sample(seq_len(n), size = 1)
```

Random Numbers in R: Recipes

Get a random integer between 1 and n

Dont:

```
sample(seq_len(n), size = 1)
```

Do:

```
sample(n, size = 1)
```

Random Numbers in R: Recipes

Get a random integer between 1 and n

Dont:

```
sample(seq_len(n), size = 1)
```



Do:

```
sample(n, size = 1)
```



`sample()` with a
scalar n samples
from 1 to n by itself

Random Numbers in R: Recipes

Get a random integer between 1 and n

Dont:

```
sample(seq_len(n), size = 1)
```



Do:

```
sample(n, size = 1)
```



or:

```
sample.int(n, size = 1)
```



`sample()` with a scalar n samples from 1 to n by itself

Random Numbers in R: Recipes

Shuffle a vector `v`

Random Numbers in R: Recipes

Shuffle a vector `v`

Dont:

```
sample(v)
```

Random Numbers in R: Recipes

Shuffle a vector `v`

Dont:

```
sample(v)
```

Do:

```
v[sample(length(v))]
```

Random Numbers in R: Recipes

Shuffle a vector `v`

Dont:

```
sample(v)
```



Do:

```
v[sample(length(v))]
```



If `v` happens to be length 1 (scalar) then `sample()` turns it into `seq_len(v)`!

Random Numbers in R: Recipes

Shuffle a vector `v`

Dont:

```
sample(v)
```



Do:

```
v[sample(length(v))]
```



If `v` happens to be length 1 (scalar) then `sample()` turns it into `seq_len(v)`!

!! Important !!

Random Numbers in R: Recipes

Shuffle a vector `v`

Dont:

```
sample(v)
```



Do:

```
v[sample(length(v))]
```



or:

```
v[sample.int(length(v))]
```



If `v` happens to be length 1 (scalar) then `sample()` turns it into `seq_len(v)`!

!! Important !!

Random Numbers in R: Recipes

Put rows of table in random order:

```
iris[sample.int(nrow(iris)), ]
```

Subsample to 10% of your data:

```
iris[sample.int(nrow(iris), size = nrow(iris) * 0.1), ]
```

Sample rows with replacement (e.g. Bootstrap):

```
iris[sample.int(nrow(iris), replace = TRUE), ]
```

Reorder a given column of a table:

```
i2 <- iris  
i2$Sepal.Width <- i2$Sepal.Width[sample.int(nrow(i2))]
```

RNGs and Reproducibility

Different RNGs in R

- Select different RNG algorithm with `RNGkind()`, or with additional arguments of `set.seed()`: `set.seed(n, kind = ...)`
- Particular properties go beyond this course (read `?Random` if interested)
- Defaults are often fine.
- Important for us:
 1. Default rng algorithms sometimes change between R versions.
 - Use `RNGversion("<R version>")` to set the PRNG `kind`, `normal.kind`, `sample.kind` to the default settings of R of version `<R version>` (e.g. `"3.6.0"`).
 - Use this if you want **reproducibility with an old R version**, but note that old R versions sometimes used sub-optimal PRNG kinds.
 2. `"L'Ecuyer-CMRG"` kind is interesting since it can be used when doing parallelization (we will see this in a later session)

PRNGs, Entropy, Security

- If I know your seed and your code, I know your random numbers!

E.g.:

[Around 1999](#), there was some online poker software that shuffled cards with an algorithm that

(1) was public and

(2) used a seed based on its system clock

- If you knew the server time approximately, and you knew a few cards, you could guess the rest of the deck

PRNGs, Entropy, Security

- If I know your seed and your code, I know your random numbers!
- **Depending on the RNG, if I know enough random values, I can deduce the RNG state and predict the following random numbers**
- E.g. If you give me three `runif()` random values, and they are
0.26550866 0.37212390 0.57285336
The next value is quite likely 0.90820779 (since I assume you used seed 1).
- E.g. "Sploosh Kaboom Probability Calculator" [[link](#)] [[video \(worth it!\)](#)]

PRNGs, Entropy, Security


- If I know your seed and your code, I know your random numbers!
 - Depending on the RNG, if I know enough random values, I can deduce the RNG state and predict the following random numbers
-
- Be careful if you generate random numbers that you don't want someone else to guess
 - Guessing your PRNG-random number is only as difficult as guessing your seed
 - so make sure the seed is *actually* random (i.e. has "[high entropy](#)") and use *cryptographically secure* PRNGs
 - As always in computer security: you better *really* know what you are doing if someone's money depends on it.

Setting a Seed

- If no seed is given, R generates one from system time and process ID
 - This is not really reproducible!
 - Always set a seed for "production" code
 - (And don't shuffle your poker decks with this)
- Seed is rounded down to an integer value, so `set.seed(1)` and `set.seed(1.8)` have the same effect.
- Seed should be between $(1-2^{31})$ and $(2^{31}-1)$ (i.e. about $\pm 2 \times 10^9$)
- When doing randomised experiments, use different seeds for different instantiations, otherwise your runs are not really independent

Setting a Seed

Different seeds -- different values?





```
> set.seed(1)
> runif(10)
[1] 0.26550866
[2] 0.37212390
[3] 0.57285336
[4] 0.90820779
[5] 0.20168193
[6] 0.89838968
[7] 0.94467527
[8] 0.66079779
[9] 0.62911404
[10] 0.06178627

> set.seed(2)
> runif(10)
[1] 0.1848823
[2] 0.7023740
[3] 0.5733263
[4] 0.1680519
[5] 0.9438393
[6] 0.9434750
[7] 0.1291590
[8] 0.8334488
[9] 0.4680185
[10] 0.5499837
```




Setting a Seed

Different seeds -- different values?

		
<code>> set.seed(0)</code>	<code>> set.seed(1)</code>	<code>> set.seed(2)</code>
<code>> runif(10)</code>	<code>> runif(10)</code>	<code>> runif(10)</code>
<div><div>[1] 0.8966972</div><div>[2] 0.2655087</div><div>[3] 0.3721239</div><div>[4] 0.5728534</div><div>[5] 0.9082078</div><div>[6] 0.2016819</div><div>[7] 0.8983897</div><div>[8] 0.9446753</div><div>[9] 0.6607978</div><div>[10] 0.6291140</div></div>	<div><div>[1] 0.26550866</div><div>[2] 0.37212390</div><div>[3] 0.57285336</div><div>[4] 0.90820779</div><div>[5] 0.20168193</div><div>[6] 0.89838968</div><div>[7] 0.94467527</div><div>[8] 0.66079779</div><div>[9] 0.62911404</div><div>[10] 0.06178627</div></div>	<div><div>[1] 0.1848823</div><div>[2] 0.7023740</div><div>[3] 0.5733263</div><div>[4] 0.1680519</div><div>[5] 0.9438393</div><div>[6] 0.9434750</div><div>[7] 0.1291590</div><div>[8] 0.8334488</div><div>[9] 0.4680185</div><div>[10] 0.5499837</div></div>

Setting a Seed

Different seeds -- different values?

<p>> set.seed(69070) > runif(10)</p>	<p>← </p>	<p>> set.seed(0) > runif(10)</p>	<p>← </p>	<p>> set.seed(1) > runif(10)</p>	<p>→ </p>	<p>> set.seed(2) > runif(10)</p>
<div data-bbox="177 535 454 1034" style="border: 1px solid red; padding: 2px;"><p>[1] 0.37212390 [2] 0.57285336 [3] 0.90820779 [4] 0.20168193 [5] 0.89838968 [6] 0.94467527 [7] 0.66079779 [8] 0.62911404 [9] 0.06178627 [10] 0.20597457</p></div>		<div data-bbox="672 633 923 1034" style="border: 1px solid red; padding: 2px;"><p>[1] 0.8966972 [2] 0.2655087 [3] 0.3721239 [4] 0.5728534 [5] 0.9082078 [6] 0.2016819 [7] 0.8983897 [8] 0.9446753 [9] 0.6607978 [10] 0.6291140</p></div>		<div data-bbox="1087 584 1348 1034" style="border: 1px solid red; padding: 2px;"><p>[1] 0.26550866 [2] 0.37212390 [3] 0.57285336 [4] 0.90820779 [5] 0.20168193 [6] 0.89838968 [7] 0.94467527 [8] 0.66079779 [9] 0.62911404 [10] 0.06178627</p></div>		<p>[1] 0.1848823 [2] 0.7023740 [3] 0.5733263 [4] 0.1680519 [5] 0.9438393 [6] 0.9434750 [7] 0.1291590 [8] 0.8334488 [9] 0.4680185 [10] 0.5499837</p>

Setting a Seed

Different seeds -- different values?

- Usually yes
- 0 and 69070 are "special" (see [here](#))
 - (The problem is: For the default RNG, `.Random.seed` is initialized with a much simpler RNG)
 - Another example: `{set.seed(9423) ; runif(2)}` and `{set.seed(27884) ; tail(runif(9), 2)}`
- Using seeds 1, 2, 3, 4, 5, ... is mostly fine if you don't use too many (<10'000s)
- Use the "L'Ecuyer-CMRG" RNG with `parallel::nextRNGStream()` if you need many independent RNG streams

Also, the streams diverge eventually!

```
> set.seed(0); round(tail(runif(229)), 3)
[1] 0.322 0.510 0.924 0.511 0.306 0.046
> set.seed(1); round(tail(runif(229)), 3)
[1] 0.510 0.924 0.511 0.258 0.046 0.418
> set.seed(69070); round(tail(runif(229)), 3)
[1] 0.924 0.511 0.258 0.701 0.418 0.854
```

RNGs, Seed, and Reproducibility

You want the same code to produce the exact same result.

- Set the same seed
- Draw random numbers in the same order
 - This breaks if you add / remove calls to `runif()` or similar between runs!
 - This breaks if you step through code manually and run `runif()` yourself sometimes!
 - This breaks if you have parts of your code that is run "optionally", e.g. may or may not use cached values!

Solutions:

- `global.seed <- 123` (or some other number) in the beginning of your script,
`set.seed(global.seed <- global.seed + 1)` after each optional code block, if you don't have too many
- Better: `set.seed(123, kind = "L'Ecuyer-CMRG")`
Then `next.seed <- parallel::nextRNGStream(.Random.seed)` before an "optional" section that may e.g. be cached
Then `set.seed(0) ; .Random.seed <- next.seed` afterwards
- [Further Reading](#)

RNGs, Seed, and Reproducibility

You want the same code to produce the exact same result.

- Set the same seed
- Draw random numbers in the same order
 - This breaks if you add / remove calls to `runif()` or similar between runs!
 - This breaks if you step through code manually and run `runif()` yourself sometimes!
 - This breaks if you have parts of your code that is run "optionally", e.g. may or may not use cached values!

Best solution: **Cache the RNG State!**

- Save RNG-state: `s <- .Random.seed`
- Restore RNG-state: `.Random.seed <- s`
- Advanced info:
 - Some RNG kind have "extra" RNG state not in `.Random.seed`. In this case:
 - Save: `s <- .Random.seed ; set.seed(1) ; .Random.seed <- s`
 - Restore: `set.seed(1) ; .Random.seed <- s`

RNGs, Seed, and Reproducibility

You want the same code to produce the exact same result.

- Set the same seed
- Draw random numbers in the same order
 - This breaks if you add / remove calls to `runif()` or similar between runs!
 - This breaks if you step through code manually and run `runif()` yourself sometimes!
 - This breaks if you have parts of your code that is run "optionally", e.g. may or may not use cached values!

Best solution: **Cache the RNG State!**

- Save RNG-state: `s <- .Random.seed`
- Restore RNG-state: `.Random.seed <- s`
- From inside functions: `.Random.seed <<- s`, since it is a *global* variable!

RNGs, Seed, and Reproducibility

You want the same code to produce the exact same result.

- Set the same seed
- Draw random numbers in the same order
 - This breaks if you add / remove calls to `runif()` or similar between runs!
 - This breaks if you step through code manually and run `runif()` yourself sometimes!
 - This breaks if you have parts of your code that is run "optionally", e.g. may or may not use cached values!
 - This can break when package versions change & packages change their behaviour

RNGs, Seed, and Reproducibility

You want the same code to produce the exact same result.

- Set the same seed
- Draw random numbers in the same order
 - This breaks if you add / remove calls to `runif()` or similar between runs!
 - This breaks if you step through code manually and run `runif()` yourself sometimes!
 - This breaks if you have parts of your code that is run "optionally", e.g. may or may not use cached values!
 - This can break when package versions change & packages change their behaviour
- Be extra careful when doing parallelization

Random Numbers: Surprises

The order in which random functions are called is really important. Because function arguments are only evaluated (or "forced") when they are first referenced, this can lead to surprises:

```
f <- function(x) {  
  y <- runif(1)  
  c(x, y)  
}
```

'x' is only forced (i.e. evaluated) *after* the call to 'runif()'

```
> set.seed(1)  
> value <- runif(1)  
> f(value)  
[1] 0.2655087 0.3721239
```

this runif(1) is evaluated *before* the function call
--> 'y <- runif(1)' was the *second* call to runif

```
> set.seed(1)  
> f(runif(1))  
[1] 0.3721239 0.2655087
```

this runif(1) is evaluated *only when 'x' is evaluated inside the function*
--> 'y <- runif(1)' was the *first* call to runif

→ Results are different, even though these two code snippets look like they should be the same.

→ It is safest to assume that *any* changes to your code change results of *all* random function calls

Random Numbers: Seed and Reproducibility II

- Sometimes other sources of randomness appear. Make sure you also set seeds for these
 - E.g. when calling an external program that uses randomness. These often have the option of setting a seed.
 - When calling python code that creates randomness, e.g. using the `reticulate` package, make sure to set the relevant seeds
 - Some old python versions [may have randomized behaviour as a security precaution](#), even if they don't do something explicitly random
 - Sometimes there are bugs in packages that make them ignore seeds, unfortunately. It is always good to check if running the script twice gives the same result (... and a true pain to debug things if this is not the case).

Random Numbers: Seed and Replications

- Sometimes you need to repeat an experiment (involving randomness) many times, so that you can do statistics with the results (estimate expected value, standard error, etc.)
- If they are in different R sessions, you need to set a different seed for all these runs.
- Using just increasing seeds > 0 is usually fine because they produce wildly different random values
- However: Try to avoid setting the *same seed* for *different code*
 - Because it could create some correlation in your results that is only due to the correlated PRNG-results
 - E.g. if your function sets a seed to `initseed` at one point, and to `initseed + 1` at another, don't call this function with values for `initseed` that differ by only 1

What We Expect You to Know

- Random value generation: `runif()`, `rexp()`, `rbinom()`, `rgeom()`, ...
- Sample and shuffle integers using `sample.int()`
- Use `sample()` on vectors if you are *sure* they are not numerics
- Use `set.seed()` to initialize the PRNG and get reproducible results
- Use `RNGversion()` to make sure the same PRNG is used in different R versions
- Get the PRNG-state from `.Random.seed`; you can save it, and then restore it by assigning a saved value to `.Random.seed`.
 - With some random functions, you have to call `set.seed(dummyvalue)` before to reset them.
 - `.Random.seed` is *not* the same value as given to `set.seed()`.
- Changing your code (or even the version of a package) can have an influence on PRNG call order (and therefore your results) in surprising ways; consider it possible that all values change once something in the code in between has changed.
- Using cached values or conditionally skipping parts of your code will lead to differing PRNG states; cache the `.Random.seed`, or call `set.seed()` again after optional chunks, to avoid the problem
- Do not run different parts of your code with the same seed.